

COMP20007 Design of Algorithms

Hashing

Daniel Beck

Lecture 15

Semester 1, 2020

Dictionaries - Recap

- Abstract Data Structure: collection of *(key, value)* pairs.

Dictionaries - Recap

- Abstract Data Structure: collection of *(key, value)* pairs.
- Required operations: Search, Insert, Delete

Dictionaries - Recap

- Abstract Data Structure: collection of *(key, value)* pairs.
- Required operations: Search, Insert, Delete
- Last lecture: Binary Search Trees (and extensions)

Dictionaries - Recap

- Abstract Data Structure: collection of *(key, value)* pairs.
- Required operations: Search, Insert, Delete
- Last lecture: Binary Search Trees (and extensions)
- This lecture: Hash Tables.

Hash Tables

- A hash table is a continuous data structure with m preallocated entries.

Hash Tables

- A hash table is a continuous data structure with m preallocated entries.
- Average case performance for Search, Insert and Delete: $\Theta(1)$

Hash Tables

- A hash table is a continuous data structure with m preallocated entries.
- Average case performance for Search, Insert and Delete: $\Theta(1)$
- Requires a *hash function*: $h(K) \rightarrow i \in [0, m - 1]$.

Hash Tables

- A hash table is a continuous data structure with m preallocated entries.
- Average case performance for Search, Insert and Delete: $\Theta(1)$
- Requires a *hash function*: $h(K) \rightarrow i \in [0, m - 1]$.
- A hash function should:
 - Be efficient ($\Theta(1)$).
 - Distribute keys evenly (uniformly) along the table.

Identity Hash Function

Question: if keys are integers, why do I need a hash function?
I could just use the key as the index, no?

Identity Hash Function

Question: if keys are integers, why do I need a hash function?
I could just use the key as the index, no?

- This is the *identity* hash function: $h(K) = K$.

Identity Hash Function

Question: if keys are integers, why do I need a hash function?
I could just use the key as the index, no?

- This is the *identity* hash function: $h(K) = K$.
- Note that $K \in [0, m - 1]$. In other words *we need to know the maximum number of keys in advance*.

Identity Hash Function

Question: if keys are integers, why do I need a hash function?
I could just use the key as the index, no?

- This is the *identity* hash function: $h(K) = K$.
- Note that $K \in [0, m - 1]$. In other words *we need to know the maximum number of keys in advance*.
- Sometimes this is possible: postcodes, for example.

Identity Hash Function

Question: if keys are integers, why do I need a hash function?
I could just use the key as the index, no?

- This is the *identity* hash function: $h(K) = K$.
- Note that $K \in [0, m - 1]$. In other words *we need to know the maximum number of keys in advance*.
- Sometimes this is possible: postcodes, for example.
- Many times it is not:
 - m is too large (need to preallocate)
 - Unbounded integers (student IDs)
 - Non-integer keys (games)

Hashing Integers

- For large/unbounded integers, an alternative function is
$$h(K) = K \bmod m$$

Hashing Integers

- For large/unbounded integers, an alternative function is
$$h(K) = K \bmod m$$
- Allow us to set the size m .

Hashing Integers

- For large/unbounded integers, an alternative function is $h(K) = K \bmod m$
- Allow us to set the size m .
- Small m results in lots of collisions, large m takes excessive memory. Best m will vary.

Hashing Strings

- Assume $A \mapsto 0$, $B \mapsto 1$, etc.
- Assume 26 characters and $m = 101$.
- Each character can be mapped to a *binary* string of length 5 ($2^5 = 32$).

Hashing Strings

- Assume $A \mapsto 0$, $B \mapsto 1$, etc.
- Assume 26 characters and $m = 101$.
- Each character can be mapped to a *binary* string of length 5 ($2^5 = 32$).

We can think of a string as a long binary number:

M Y K E Y \mapsto 0110011000010100010011000 (= 13379736)

$$13379736 \bmod 101 = 64$$

So 64 is the position of string M Y K E Y in the hash table.

Hashing Strings

We deliberately chose m to be **prime**.

$$13379736 = 12 \times 32^4 + 24 \times 32^3 + 10 \times 32^2 + 4 \times 32 + 24$$

With $m = 32$, the hash value of any key is the last character's value!

Hashing Long Strings

Assume *chr* be the function that gives a character's number, so for example, $chr(c) = 2$.

Hashing Long Strings

Assume chr be the function that gives a character's number, so for example, $chr(c) = 2$.

Then we have

$$h(s) = (\sum_{i=0}^{|s|-1} chr(s_i) \times 32^{|s|-i-1}) \bmod m,$$

where m is a prime number. For example,

$$h(\text{V E R Y L O N G K E Y}) = (21 \times 32^{10} + 4 \times 32^9 + \dots) \bmod 101$$

Hashing Long Strings

Assume chr be the function that gives a character's number, so for example, $chr(c) = 2$.

Then we have

$$h(s) = (\sum_{i=0}^{|s|-1} chr(s_i) \times 32^{|s|-i-1}) \bmod m,$$

where m is a prime number. For example,

$$h(\text{V E R Y L O N G K E Y}) = (21 \times 32^{10} + 4 \times 32^9 + \dots) \bmod 101$$

The term between parenthesis can become quite large and result in overflow.

Horner's Rule

Instead of

$$21 \times 32^{10} + 4 \times 32^9 + 17 \times 32^8 + 24 \times 32^7 \dots$$

factor out repeatedly:

$$(\dots ((21 \times 32 + 4) \times 32 + 17) \times 32 + \dots) + 24$$

Horner's Rule

Instead of

$$21 \times 32^{10} + 4 \times 32^9 + 17 \times 32^8 + 24 \times 32^7 \dots$$

factor out repeatedly:

$$(\dots ((21 \times 32 + 4) \times 32 + 17) \times 32 + \dots) + 24$$

Now utilize these properties of modular arithmetic:

$$(x + y) \bmod m = ((x \bmod m) + (y \bmod m)) \bmod m$$

$$(x \times y) \bmod m = ((x \bmod m) \times (y \bmod m)) \bmod m$$

So for each sub-expression it suffices to take values modulo m .

Collisions

Happens when the hash function give identical results to two different keys.

Collisions

Happens when the hash function give identical results to two different keys.

We saw two solutions:

- Separate Chaining
- Linear Probing

Collisions

Happens when the hash function give identical results to two different keys.

We saw two solutions:

- Separate Chaining
- Linear Probing

Practical efficiency will depend on the table **load factor**:

$$\alpha = n/m$$

Separate Chaining

Assign multiple records per cell (usually through a linked list)

Separate Chaining

Assign multiple records per cell (usually through a linked list)

- Assuming even distribution of the n keys.

Separate Chaining

Assign multiple records per cell (usually through a linked list)

- Assuming even distribution of the n keys.
- A successful search requires $1 + \alpha/2$ operations on average.

Separate Chaining

Assign multiple records per cell (usually through a linked list)

- Assuming even distribution of the n keys.
- A successful search requires $1 + \alpha/2$ operations on average.
- An unsuccessful search requires α operations on average.

Separate Chaining

Assign multiple records per cell (usually through a linked list)

- Assuming even distribution of the n keys.
- A successful search requires $1 + \alpha/2$ operations on average.
- An unsuccessful search requires α operations on average.
- Almost same numbers for Insert and Delete.

Separate Chaining

Assign multiple records per cell (usually through a linked list)

- Assuming even distribution of the n keys.
- A successful search requires $1 + \alpha/2$ operations on average.
- An unsuccessful search requires α operations on average.
- Almost same numbers for Insert and Delete.
- Worst case $\Theta(n)$ only with a bad hash function (load factor is more of an issue).

Separate Chaining

Assign multiple records per cell (usually through a linked list)

- Assuming even distribution of the n keys.
- A successful search requires $1 + \alpha/2$ operations on average.
- An unsuccessful search requires α operations on average.
- Almost same numbers for Insert and Delete.
- Worst case $\Theta(n)$ only with a bad hash function (load factor is more of an issue).
- Requires extra memory.

Linear Probing

Populate successive empty cells.

Linear Probing

Populate successive empty cells.

- Much harder analysis, simplified results show:
- A successful search requires $(1/2) \times (1 + 1/(1 - \alpha))$ operations on average.
- An unsuccessful search requires $(1/2) \times (1 + 1/(1 - \alpha)^2)$ operations on average.

Linear Probing

Populate successive empty cells.

- Much harder analysis, simplified results show:
- A successful search requires $(1/2) \times (1 + 1/(1 - \alpha))$ operations on average.
- An unsuccessful search requires $(1/2) \times (1 + 1/(1 - \alpha)^2)$ operations on average.
- Similar numbers for Insert. Delete virtually impossible.

Linear Probing

Populate successive empty cells.

- Much harder analysis, simplified results show:
- A successful search requires $(1/2) \times (1 + 1/(1 - \alpha))$ operations on average.
- An unsuccessful search requires $(1/2) \times (1 + 1/(1 - \alpha)^2)$ operations on average.
- Similar numbers for Insert. Delete virtually impossible.
- Does not require extra memory.

Linear Probing

Populate successive empty cells.

- Much harder analysis, simplified results show:
- A successful search requires $(1/2) \times (1 + 1/(1 - \alpha))$ operations on average.
- An unsuccessful search requires $(1/2) \times (1 + 1/(1 - \alpha)^2)$ operations on average.
- Similar numbers for Insert. Delete virtually impossible.
- Does not require extra memory.
- Worst case $\Theta(n)$ with a bad hash function *and/or* clusters.

Double Hashing

A generalisation of Linear Probing.

Apply a second hash function in case of collision.

Double Hashing

A generalisation of Linear Probing.

Apply a second hash function in case of collision.

- First try: $h(K)$
- Second try: $(h(K) + s(K)) \bmod m$
- Third try: $(h(K) + 2s(K)) \bmod m$
- ...

Double Hashing

A generalisation of Linear Probing.

Apply a second hash function in case of collision.

- First try: $h(K)$
- Second try: $(h(K) + s(K)) \bmod m$
- Third try: $(h(K) + 2s(K)) \bmod m$
- ...

Another reason to use prime m in $h(K)$: will guarantee to find a free cell if there is one.

Double Hashing

A generalisation of Linear Probing.

Apply a second hash function in case of collision.

- First try: $h(K)$
- Second try: $(h(K) + s(K)) \bmod m$
- Third try: $(h(K) + 2s(K)) \bmod m$
- ...

Another reason to use prime m in $h(K)$: will guarantee to find a free cell if there is one.

Both Linear Probing and Double Hashing are sometimes referred as *Open Addressing* methods.

Rehashing

- High load factors deteriorate the performance of a hash table (for linear probing, ideally we should have $\alpha < 0.9$).

Rehashing

- High load factors deteriorate the performance of a hash table (for linear probing, ideally we should have $\alpha < 0.9$).
- *Rehashing* allocates a new table (usually around double the size) and move every item from the previous table to the new one.

Rehashing

- High load factors deteriorate the performance of a hash table (for linear probing, ideally we should have $\alpha < 0.9$).
- *Rehashing* allocates a new table (usually around double the size) and move every item from the previous table to the new one.
- *Very expensive operation*, but happens infrequently.

Summary

Hash Tables:

Summary

Hash Tables:

- Implement dictionaries.

Summary

Hash Tables:

- Implement dictionaries.
- Allow $\Theta(1)$ Search, Insert and Delete in the average case.

Summary

Hash Tables:

- Implement dictionaries.
- Allow $\Theta(1)$ Search, Insert and Delete in the average case.
- Preallocates memory (size m).

Summary

Hash Tables:

- Implement dictionaries.
- Allow $\Theta(1)$ Search, Insert and Delete in the average case.
- Preallocates memory (size m).
- Requires good *hash functions*.

Summary

Hash Tables:

- Implement dictionaries.
- Allow $\Theta(1)$ Search, Insert and Delete in the average case.
- Preallocates memory (size m).
- Requires good *hash functions*.
- Requires good collision handling.

Pros and Cons

If Hash Tables are so good, why bother with BSTs?

Pros and Cons

If Hash Tables are so good, why bother with BSTs?

- Hash Tables ignore *key ordering*, unlike BSTs.
- Queries like “give me all records with keys between 100 and 200” are easy within a BST but much less efficient in a hash table.

Pros and Cons

If Hash Tables are so good, why bother with BSTs?

- Hash Tables ignore *key ordering*, unlike BSTs.
- Queries like “give me all records with keys between 100 and 200” are easy within a BST but much less efficient in a hash table.
- Also: memory requirements of a hash table are much higher.

Pros and Cons

If Hash Tables are so good, why bother with BSTs?

- Hash Tables ignore *key ordering*, unlike BSTs.
- Queries like “give me all records with keys between 100 and 200” are easy within a BST but much less efficient in a hash table.
- Also: memory requirements of a hash table are much higher.

That being said, if hashing is applicable, a well-tuned hash table will typically outperform BSTs.

Python dictionaries (*dict* type)

Python dictionaries (*dict* type)

- Open addressing using *pseudo-random probing*
- Rehashing happens when $\alpha = 2/3$

Python dictionaries (*dict* type)

- Open addressing using *pseudo-random probing*
- Rehashing happens when $\alpha = 2/3$

C++ *unordered_maps*

In Practice

Python dictionaries (*dict* type)

- Open addressing using *pseudo-random probing*
- Rehashing happens when $\alpha = 2/3$

C++ *unordered_maps*

- Uses chaining.
- Rehashing happens when $\alpha = 1$

In Practice

Python dictionaries (*dict* type)

- Open addressing using *pseudo-random probing*
- Rehashing happens when $\alpha = 2/3$

C++ *unordered_maps*

- Uses chaining.
- Rehashing happens when $\alpha = 1$

Next lecture: what happens if records/data is too large?